Storing and Accessing Data row

Lesson 02: Primary Index Mechanics

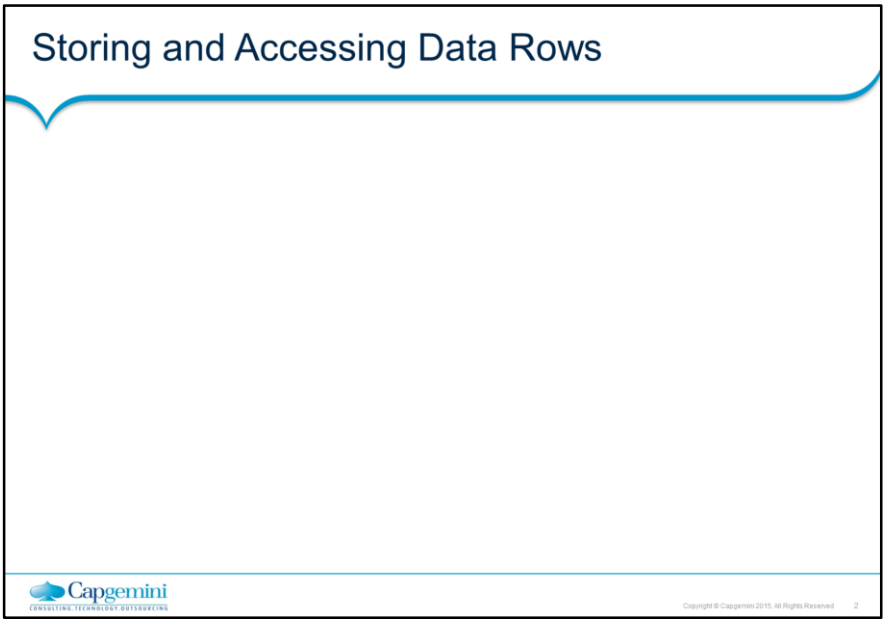
Storing and Accessing Data Rows
Capgemini
CONSULTING. TECHNOLOGY. OUTSOURCING
Copyright © Capgemini 2015. All Rights Reserved 2

How Does Teradata Store Rows?

- Teradata uses hash partitioning and distribution to randomly and evenly distribute data
- across all AMPs.
- The rows of every table are distributed among all AMPs - and ideally will be evenly distributed among all AMPs.
- Each AMP is responsible for a subset of the rows of each table. Evenly distributed tables result in evenly distributed workloads. The data is not placed in any particular order

The benefits of unordered data include:
- No maintenance needed to preserve order, and
- It is independent of any query being submitted.

The benefits of automatic data placement include:
- Distribution is the same regardless of data
- Distribution is based on row content, not data

Capgemini
CONSULTING. TECHNOLOGY. OUTSOURCING

How do other databases Store Rows

- Many use range distribution Creates intensive maintenance for DBA

DBA's must consider:

- How to partition the data
- How large to make the partitions Where is there data contention How are users accessing the data

Placing all data into a single partition creates bottlenecks for all queries against that data.

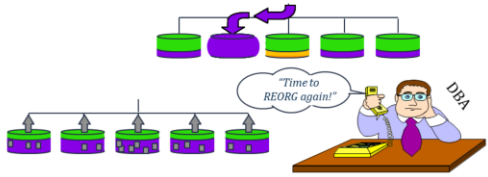

Teradata DBAs never need to do costly reorganizations!

Capgemini
CONSULTING. TECHNOLOGY. OUTSOURCING

Primary Indexes

- The mechanism used to assign a row to an AMP
- A table must have a Primary Index
- The Primary Index cannot be changed

UPI UPI's guarantee even data distribution and eliminate duplicate row checking.

- If the index choice of column(s) is unique, we call this a *UPI* (Unique Primary Index).
- A UPI choice will result in even distribution of the rows of the table across all AMPs.

NUPI

- If the index choice of column(s) isn't unique, we call this a *NUPI* (Non-Unique Primary Index).
- A NUPI choice will result in even distribution of the rows of the table proportional to the degree of uniqueness of the index.

Capgemini
CONSULTING. TECHNOLOGY. OUTSOURCING

Creating a Primary Index

- A Primary Index is defined at table creation.
- It may consist of a single column or a combination of up to 16 columns.

UPI

```
CREATE TABLE sample_1
            (col_a      INT
            ,col_b      INT
            ,col_c      INT)
UNIQUE PRIMARY INDEX (col_b);
```

NUPI

```
CREATE TABLE sample_2
            (col_x      INT
            ,col_y      INT
            ,col_z      INT)
PRIMARY INDEX (col_x);
```

Note: Changing the Primary Index requires dropping and recreating the table.

Capgemini
CONSULTING. TECHNOLOGY. OUTSOURCING

Primary Index Values

- The value of the Primary Index for a specific row determines its AMP assignment.
- This is done using the hashing algorithm.

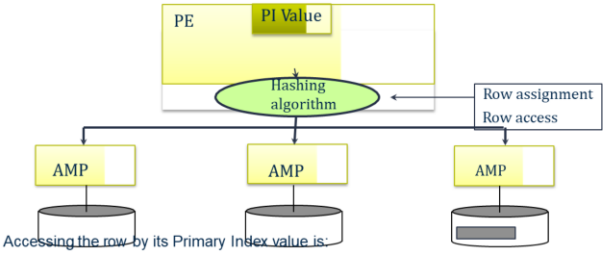

Accessing the row by its Primary Index value is:

- Always a *one-AMP* operation
- The most efficient way to access a row

Other table access techniques:
- Secondary index access
- Full table scans

Capgemini
CONSULTING. TECHNOLOGY. OUTSOURCING

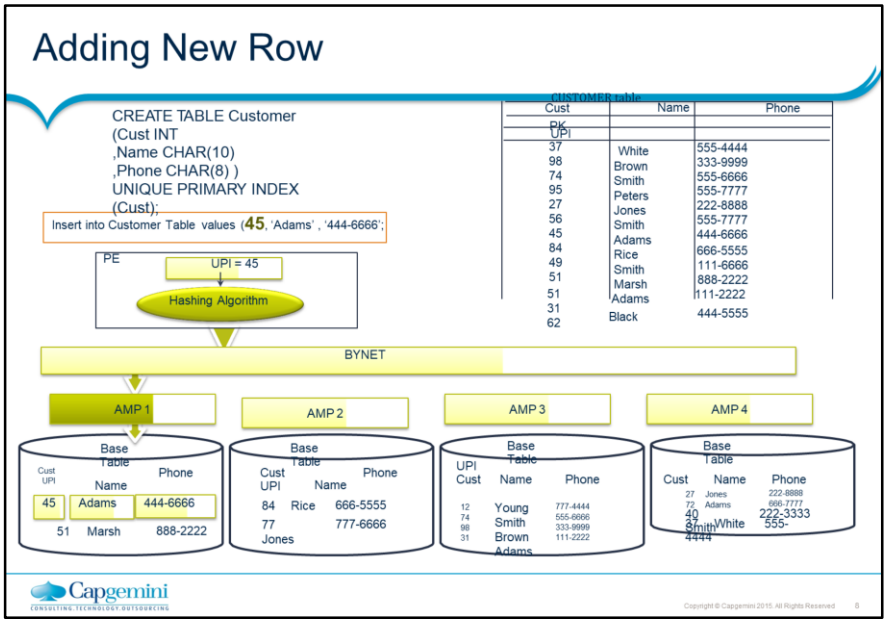
Adding New Row
CREATE TABLE Customer
(Cust INT
,Name CHAR(10)
,Phone CHAR(8))
UNIQUE PRIMARY INDEX
(Cust);
Insert into Customer Table values (45, 'Adams' , '444-6666';
PE
UPI = 45
Hashing Algorithm
BYNET
AMP 1
AMP 2
AMP 3
AMP 4
Base Table
Cust UPI
Name
Phone
45
Adams
444-6666
51 Marsh 888-2222
84 Rice 666-5555
77 Jones 777-6666
12 Young 777-4444
74 Smith 555-6666
98 Brown 333-9999
31 Adams 111-2222
27 Jones 222-8888
72 Adams 666-7777
40 White 222-3333
CUSTOMER table
Cust
Name
Phone
PK
UPI
37 White 555-4444
98 Brown 333-9999
74 Smith 555-6666
95 Peters 555-7777
27 Jones 222-8888
56 Smith 555-7777
45 Adams 444-6666
84 Rice 666-5555
49 Smith 111-6666
51 Marsh 888-2222
51 Adams 111-2222
31 Black 444-5555
62
Capgemini
CONSULTING. TECHNOLOGY. OUTSOURCING
Copyright © Capgemini 2015. All Rights Reserved 8

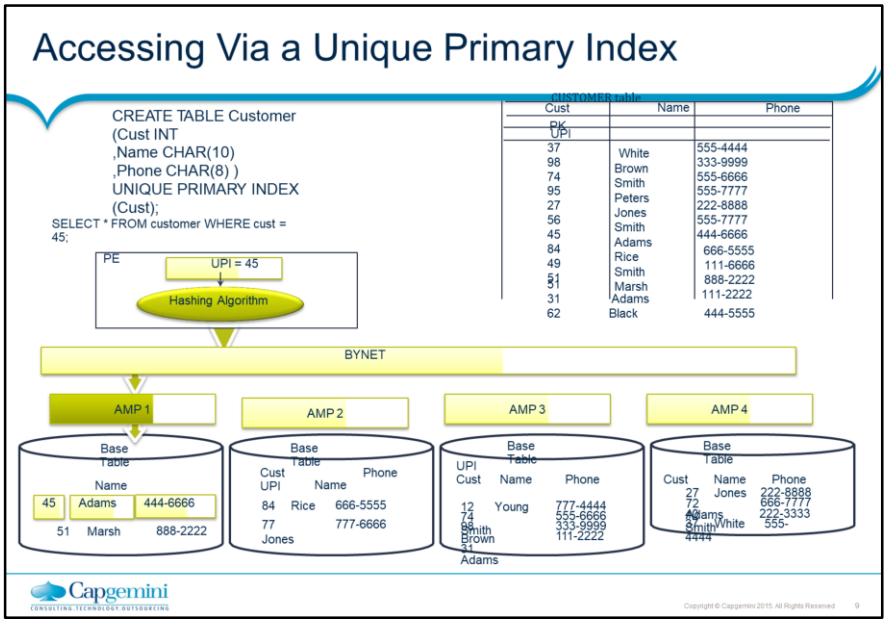
Accessing Via a Unique Primary Index
CREATE TABLE Customer
(Cust INT
,Name CHAR(10)
,Phone CHAR(8))
UNIQUE PRIMARY INDEX
(Cust);
SELECT * FROM customer WHERE cust = 45;
PE
UPI = 45
Hashing Algorithm
BYNET
AMP 1
AMP 2
AMP 3
AMP 4
Base Table
Name
45 Adams 444-6666
51 Marsh 888-2222
Cust UPI Name Phone
84 Rice 666-5555
77 Jones 777-6666
UPI Cust Name Phone
12 Young 777-4444
74 555-6666
96 333-9999
Brown 111-2222
31 Adams
Cust Name Phone
27 Jones 222-8888
72 666-7777
222-3333
White 555-
4444
CUSTOMER table
Cust Name Phone
PK
UPI
37 555-4444
98 White 333-9999
74 Brown 555-6666
95 Smith 555-7777
27 Peters 222-8888
56 Jones 555-7777
45 Smith 444-6666
84 Adams 666-5555
49 Rice 111-6666
51 Smith 888-2222
31 Marsh 111-2222
62 Adams
Black 444-5555
Capgemini
CONSULTING. TECHNOLOGY. OUTSOURCING
Copyright © Capgemini 2015. All Rights Reserved 9

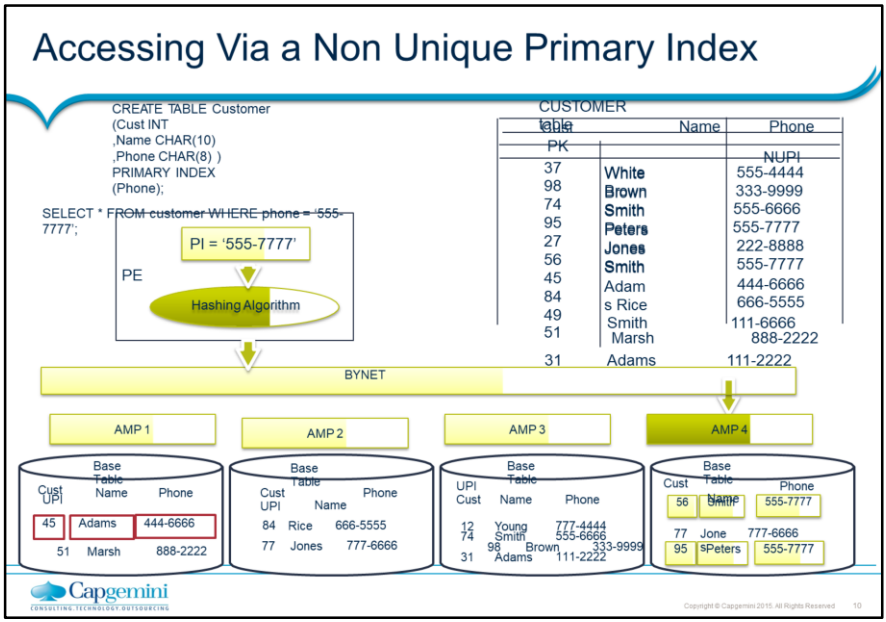
Accessing Via a Non Unique Primary Index
CREATE TABLE Customer
(Cust INT
,Name CHAR(10)
,Phone CHAR(8))
PRIMARY INDEX
(Phone);
SELECT * FROM customer WHERE phone = '555-7777';
PI = '555-7777'
PE
Hashing Algorithm
BYNET
CUSTOMER table
Cust PK
Name
Phone NUPI
37 White 555-4444
98 Brown 333-9999
74 Smith 555-6666
95 Peters 555-7777
27 Jones 222-8888
56 Smith 555-7777
45 Adam 444-6666
84 s Rice 666-5555
49 Smith 111-6666
51 Marsh 888-2222
31 Adams 111-2222
AMP 1
AMP 2
AMP 3
AMP 4
Base Table
Cust UPI Name Phone
45 Adams 444-6666
51 Marsh 888-2222
Base Table
Cust UPI Name Phone
84 Rice 666-5555
77 Jones 777-6666
Base Table
UPI Cust Name Phone
12 Young 777-4444
74 Smith 555-6666
98 Brown 333-9999
31 Adams 111-2222
Base Table
Cust Name Phone
56 Smith 555-7777
77 Jone 777-6666
95 sPeters 555-7777
Capgemini
CONSULTING. TECHNOLOGY. OUTSOURCING
Copyright © Capgemini 2015. All Rights Reserved 10

Primary Keys and Primary Indexes

- *Indexes* are conceptually different from *keys:*
  - A *PK* is a relational modeling convention which allows each row to be uniquely identified.
  - A *PI* is a Teradata convention which determines how the row will be stored and accessed.

Primary Key	Primary Index
Logical concept of data modeling	Physical mechanism for access and storage
Teradata doesn't need to recognize	Each table must have exactly one
No limit on column numbers	16-column limit
Documented in data model (Optional in CREATE TABLE)	Defined in CREATE TABLE statement
Must be unique	May be unique or non-unique
Uniquely identifies each row	Used to place and locate each row on an AMP
Values should not change	Values may be changed (Del+ Ins)
May not be NULL—requires a value	May be NULL
Does not imply an access path	Defines most efficient access path
Chosen for logical correctness	Chosen for physical performance

- A significant percentage of tables may use the same columns for both the PK and the PI.
- A well-designed database will use a PI that is different from the PK for some tables.

Capgemini
CONSULTING. TECHNOLOGY. OUTSOURCING

 11

Row Distribution Using an UPI

Order

Order Number	Customer Number	Order Date	Order Status
PK			
UPI			
7325	2	4/13	O
7324	3	4/13	O
7415	1	4/13	C
7103	1	4/10	O
7225	2	4/15	C
7384	1	4/12	C
7402	3	4/16	C
7188	1	4/13	C
7202	2	4/09	C

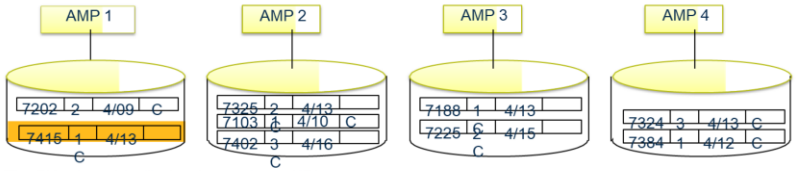

The PK column(s) will often be used as a UPI.
PI values for Order_Number are known to be unique (it's a PK).

- Teradata will distribute different index values evenly across all AMPs.

Resulting row distribution among AMPs is uniform.

Capgemini
CONSULTING. TECHNOLOGY. OUTSOURCING

 12

Row Distribution Using an NUPI

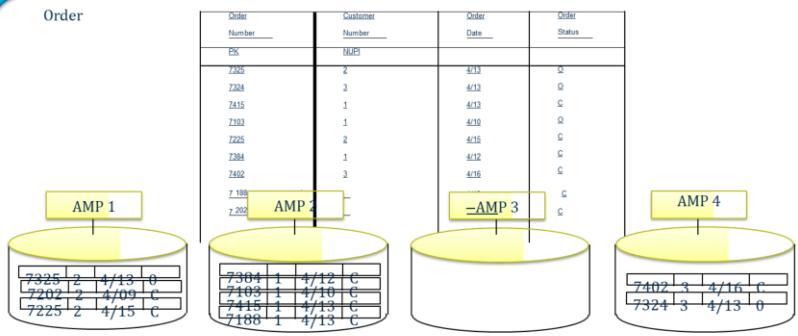

- Customer_Number may be the preferred access column for ORDER table, thus a good index candidate.
- Values for Customer_Number are non-unique and therefore a NUPI.
- Rows with the same PI value distribute to the same AMP causing row distribution to be less uniform or skewed.

Row Distribution Using a Highly Non Unique Index

Order

Order Number	Customer Number	Order Date	Order Status
PK			
			NUPI
7325	2	4/13	O
7324	3	4/13	O
7415	1	4/13	C
7103	1	4/10	C
7225	2	4/15	C
7384	1	4/12	
7402	3	4/16	
7188	1	4/13	
7202	2	4/09	

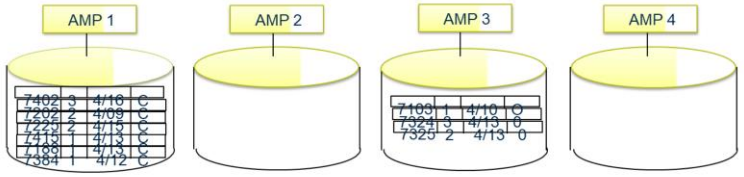

- Values for Order_Status are highly non-unique, and therefore, it is a NUPI.
- Only two values exist, so only two AMPs will ever be used for this table.
- This table will not perform well in parallel operations.
- Highly non-unique columns are poor PI choices.
- The degree of uniqueness is critical to efficiency.

Capgemini
CONSULTING. TECHNOLOGY. OUTSOURCING

 14

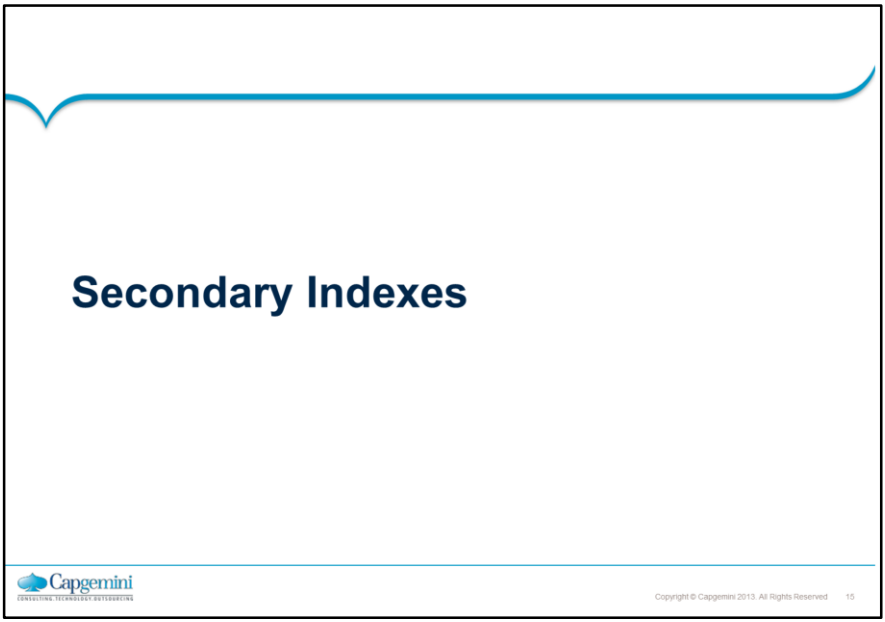
Secondary Indexes
Capgemini
CONSULTING. TECHNOLOGY. OUTSOURCING
Copyright © Capgemini 2013. All Rights Reserved 15

Secondary Indexes

- A secondary index is an alternate path to the rows of a table.
- A table can have from 0 to 32 secondary indexes. Secondary indexes:
  - Do not affect table distribution.
  - Add overhead, both in terms of disk space and maintenance.
  - May be added or dropped dynamically as needed. Are chosen to improve table performance.

There are three general ways to access a table:

- Primary index access (*one*-AMP access)
- Secondary index access (*two*-or *all*-AMP access)
- Full Table Scan (*all*-AMP access)

Capgemini
CONSULTING. TECHNOLOGY. OUTSOURCING

Choosing a Secondary Index

A secondary index may be defined:

->At table creation (CREATE TABLE)

-> Following table creation (CREATE NDEX)

-> Using up to 16 columns

USI

- If the index choice of column(s) is unique, it is called a USI (unique secondary index).
- Accessing a row via a USI typically requires 2 AMPs.

NUS

I

- If the index choice of column(s) is non-unique, it is called a NUSI (non-unique secondary index).
- Accessing a row via a NUSI requires all AMPs.

USI

```
CREATE UNIQUE
INDEX
(employee-number) n employee
```

NUS

I

```
CREATE INDEX (last-name) on
employee
CREATE INDEX (last-name, first-name) on
employee
```

Note:

- Secondary indexes cause an internal sub-table to be built.
- Dropping the index causes the sub-table to be deleted.

Capgemini
CONSULTING. TECHNOLOGY. OUTSOURCING

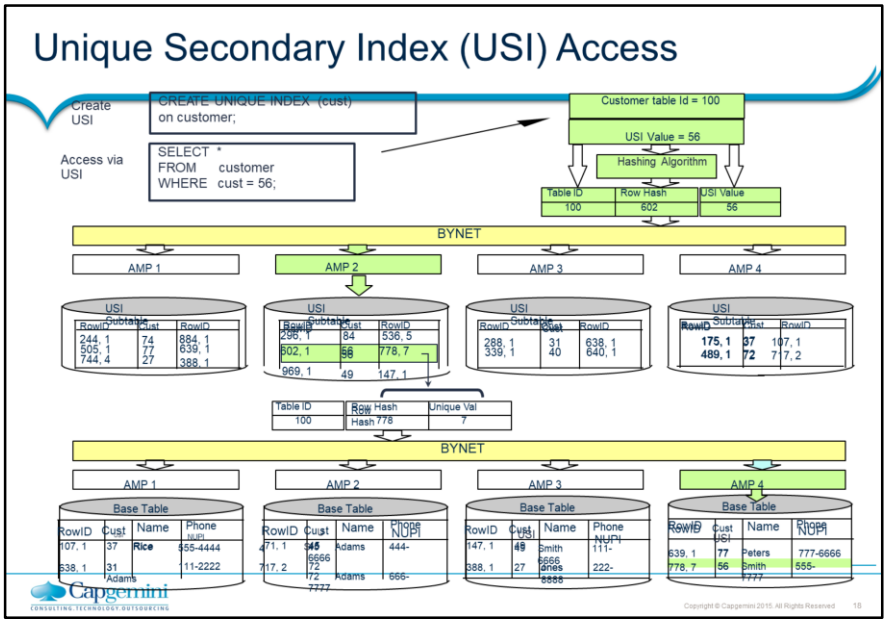
Unique Secondary Index (USI) Access
Create USI
CREATE UNIQUE INDEX (cust)
on customer;
Access via USI
SELECT *
FROM customer
WHERE cust = 56;
Customer table Id = 100
USI Value = 56
Hashing Algorithm
Table ID
Row Hash
USI Value
100
602
56
BYNET
AMP 1
AMP 2
AMP 3
AMP 4
USI
Subtable
RowID
Cust
244, 1
505, 1
744, 4
74
77
27
884, 1
639, 1
388, 1
296, 1
602, 1
969, 1
84
56
49
536, 5
778, 7
147, 1
288, 1
339, 1
31
40
638, 1
640, 1
175, 1
489, 1
37
72
107, 1
777, 2
Table ID
Row Hash
Unique Val
100
Hash 778
7
Base Table
Name
Phone
NUPI
107, 1
37
Rice
555-4444
638, 1
31
Adams
111-2222
471, 1
717, 2
Adams
444-
666-
147, 1
388, 1
Smith
222-
639, 1
778, 7
Peters
Smith
777-6666
555-
Capgemini
CONSULTING. TECHNOLOGY. OUTSOURCING
Copyright © Capgemini 2015. All Rights Reserved
18

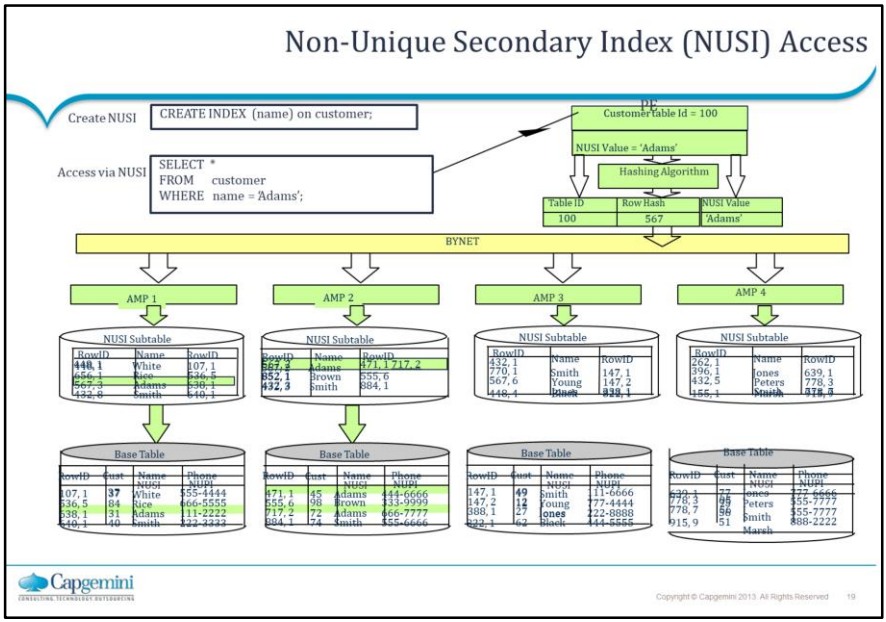
Non-Unique Secondary Index (NUSI) Access
Create NUSI
CREATE INDEX (name) on customer;
Access via NUSI
SELECT *
FROM customer
WHERE name = 'Adams';
PE
Customer Table Id = 100
NUSI Value = 'Adams'
Hashing Algorithm
Table ID
Row Hash
NUSI Value
100
567
'Adams'
BYNET
AMP 1
AMP 2
AMP 3
AMP 4
NUSI Subtable
RowID
Name
RowID
Base Table
RowID
Cust
Name
Phone
NUPI
NUSI
Capgemini
Copyright © Capgemini 2013. All Rights Reserved
19

Comparison of Primary and Secondary Indexes

Index Feature	Primary	Secondary
Required?	Yes	No
Number per Table	1	0-32
Max Number of Columns	16	16
Unique or Non-Unique?	Both	Both
Affects Row Distribution	Yes	No
Created/Dropped Dynamically	No	Yes
Improves Access	Yes	Yes
Multiple Data Types	Yes	Yes
Separate Physical Structure	None	Sub-table
Extra Processing Overhead	No	Yes

Capgemini
CONSULTING. TECHNOLOGY. OUTSOURCING

Full – Table Scans

- Every row of the table must be read.
- All AMPs scan their portion of the table in parallel. Primary Index choice affects FTS performance.
- Full-table scans typically occur when either:
  - The index columns are not used in the query
  - An index is used in a non-equality test
  - A range of values is specified for the primary index

CUSTOMER

Cust_ID	Cust_Name	Cust_Phone
USI	NUSI	NUPI

Examples of Full-Table Scans:

```
SELECT  *  FROM customer WHERE Cust_Phone LIKE '524-_ _ _ _';
SELECT  * FROM customer WHERE Cust_Name <> 'Davis';

SELECT  * FROM customer WHERE Cust_ID > 1000;
```

Capgemini
CONSULTING. TECHNOLOGY. OUTSOURCING

